

**UM6P — Mohammed VI Polytechnic University**

College of Computing

# **Simulation of an Artificial Neural Network (MLP)**

Parallelization using OpenMP and MPI

**High Performance Computing Project**

**Prepared by:**

Oumkalthoum M'HAMDI  
Zineb MIFTAH

**Academic Year: 2024–2025**

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Project Context . . . . .	1
1.2	Objectives . . . . .	1
1.3	Report Organization . . . . .	2
<b>2</b>	<b>Theoretical Background</b>	<b>3</b>
2.1	Neural Network Fundamentals . . . . .	3
2.1.1	Multi-Layer Perceptron Architecture . . . . .	3
2.1.2	Activation Functions . . . . .	4
2.1.3	Softmax and Classification . . . . .	4
2.2	Loss Function and Regularization . . . . .	4
2.2.1	Cross-Entropy Loss . . . . .	4
2.2.2	L2 Regularization . . . . .	4
2.3	Batch Gradient Descent . . . . .	5
2.3.1	Full-Batch Gradient Descent . . . . .	5
2.3.2	Mini-Batch Gradient Descent . . . . .	5
2.4	Learning Rate Scheduling . . . . .	5
2.4.1	Fixed Learning Rate . . . . .	6
2.4.2	Exponential Decay . . . . .	6
2.4.3	Inverse Time Decay . . . . .	6
<b>3</b>	<b>Memory Management and Profiling</b>	<b>7</b>
3.1	Valgrind Memcheck . . . . .	7
3.1.1	Procedure . . . . .	7
3.1.2	Before Optimization . . . . .	7
3.1.3	After Optimization . . . . .	8
3.2	Callgrind Profiling . . . . .	8
3.2.1	Procedure . . . . .	8
3.2.2	Report . . . . .	8
3.3	Interpretation . . . . .	9
3.3.1	Key Findings . . . . .	9
<b>4</b>	<b>Training Optimization</b>	<b>10</b>
4.1	Procedure . . . . .	10
4.2	Mini-Batch Gradient Descent . . . . .	10
4.2.1	Implementation Details . . . . .	10
4.2.2	Before: Full-Batch . . . . .	11

4.2.3	After: Mini-Batch . . . . .	11
4.3	Activation Function Optimization . . . . .	12
4.3.1	Benchmark Setup . . . . .	12
4.3.2	Results . . . . .	12
4.3.3	Analysis . . . . .	12
4.4	Learning Rate Scheduling . . . . .	13
4.4.1	Benchmark Setup . . . . .	13
4.4.2	Results . . . . .	13
4.4.3	Analysis . . . . .	13
<b>5</b>	<b>Parallelization</b>	<b>15</b>
5.1	OpenMP Parallelization . . . . .	15
5.1.1	Procedure . . . . .	15
5.1.2	Baseline (Before Parallelization) . . . . .	16
5.2	MPI/Hybrid Parallelization . . . . .	17
<b>6</b>	<b>Conclusion</b>	<b>20</b>
6.1	Summary of Achievements . . . . .	20
6.2	Key Findings . . . . .	20
6.2.1	Algorithmic Insights . . . . .	20
6.2.2	Parallel Computing Lessons . . . . .	20
6.3	Optimal Configuration . . . . .	21
6.4	Lessons Learned . . . . .	21
6.5	Conclusion . . . . .	21
.1	Code Snippets . . . . .	23
.1.1	Matrix Multiplication with OpenMP . . . . .	23
.1.2	MPI Gradient Aggregation . . . . .	23
.2	Compilation Instructions . . . . .	24
.2.1	OpenMP Version . . . . .	24
.2.2	MPI/Hybrid Version . . . . .	24

# Chapter 1

## Introduction

Neural networks have become fundamental tools in machine learning, but training them efficiently remains computationally challenging. This project aims to improve the performance, efficiency, and scalability of a C implementation of a simple feedforward neural network using modern high-performance computing techniques.

### 1.1 Project Context

The original implementation uses full-batch gradient descent with static learning rates and manual memory management without optimization. This baseline provides significant opportunities for performance improvements through:

- Memory leak detection and correction
- Computational hotspot identification
- Training algorithm optimization
- Parallel computing with OpenMP and MPI

### 1.2 Objectives

The primary objectives of this project are:

1. **Memory Safety:** Eliminate memory leaks and ensure proper resource management
2. **Performance Profiling:** Identify computational bottlenecks using Valgrind tools
3. **Training Optimization:** Implement mini-batch gradient descent and learning rate scheduling
4. **Activation Functions:** Evaluate different activation functions (ReLU, Sigmoid, Leaky ReLU, Tanh)
5. **Parallelization:** Achieve scalable performance using OpenMP and MPI

## 1.3 Report Organization

This report is structured as follows: Section 2 introduces the theoretical foundations of neural networks. Section 3 presents memory management and profiling results. Section 4 describes training optimizations and benchmarks. Section 5 details parallelization strategies and performance analysis. Finally, Section 6 concludes with key findings.

# Chapter 2

## Theoretical Background

This section introduces the fundamental concepts underlying our neural network implementation.

### 2.1 Neural Network Fundamentals

#### 2.1.1 Multi-Layer Perceptron Architecture

A Multi-Layer Perceptron (MLP) is a feedforward neural network consisting of:

- **Input Layer:**  $d$  input features  $\mathbf{x} \in \mathbb{R}^d$
- **Hidden Layer:**  $h$  hidden units with activation function  $f(\cdot)$
- **Output Layer:**  $c$  output classes with softmax activation

The forward propagation is defined as:

$$\mathbf{z}^{(1)} = \mathbf{x}\mathbf{W}^{(1)} + \mathbf{b}^{(1)} \quad (2.1)$$

$$\mathbf{a}^{(1)} = f(\mathbf{z}^{(1)}) \quad (2.2)$$

$$\mathbf{z}^{(2)} = \mathbf{a}^{(1)}\mathbf{W}^{(2)} + \mathbf{b}^{(2)} \quad (2.3)$$

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{z}^{(2)}) \quad (2.4)$$

where:

- $\mathbf{W}^{(1)} \in \mathbb{R}^{d \times h}$ : Input-to-hidden weights
- $\mathbf{b}^{(1)} \in \mathbb{R}^h$ : Hidden layer biases
- $\mathbf{W}^{(2)} \in \mathbb{R}^{h \times c}$ : Hidden-to-output weights
- $\mathbf{b}^{(2)} \in \mathbb{R}^c$ : Output layer biases

[Image of a two layer neural network architecture (MLP)]

### 2.1.2 Activation Functions

We evaluate four activation functions in this project:

Function	Formula	Derivative
Tanh	$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$	$1 - \tanh^2(x)$
ReLU	$\max(0, x)$	$\begin{cases} 1 & x > 0 \\ 0 & x \leq 0 \end{cases}$
Leaky ReLU	$\max(0.01x, x)$	$\begin{cases} 1 & x > 0 \\ 0.01 & x \leq 0 \end{cases}$
Sigmoid	$\frac{1}{1+e^{-x}}$	$\sigma(x)(1 - \sigma(x))$

Table 2.1: Activation functions and their derivatives

### 2.1.3 Softmax and Classification

The softmax function converts raw scores into probabilities:

$$\text{softmax}(\mathbf{z})_i = \frac{e^{z_i - \max(\mathbf{z})}}{\sum_{j=1}^c e^{z_j - \max(\mathbf{z})}} \quad (2.5)$$

The  $\max(\mathbf{z})$  subtraction ensures numerical stability.

## 2.2 Loss Function and Regularization

### 2.2.1 Cross-Entropy Loss

For classification, we use the cross-entropy loss:

$$\mathcal{L}_{\text{CE}} = -\frac{1}{N} \sum_{i=1}^N \log(\hat{y}_i^{(y_i)}) \quad (2.6)$$

where  $\hat{y}_i^{(y_i)}$  is the predicted probability for the true class  $y_i$  of sample  $i$ .

### 2.2.2 L2 Regularization

To prevent overfitting, we add L2 regularization:

$$\mathcal{L}_{\text{reg}} = \frac{\lambda}{2} (\|\mathbf{W}^{(1)}\|_F^2 + \|\mathbf{W}^{(2)}\|_F^2) \quad (2.7)$$

The total loss becomes:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{CE}} + \mathcal{L}_{\text{reg}} \quad (2.8)$$

## 2.3 Batch Gradient Descent

### 2.3.1 Full-Batch Gradient Descent

The original implementation uses full-batch gradient descent:

---

**Algorithm 1** Full-Batch Gradient Descent

---

```

1: Initialize weights  $\mathbf{W}^{(1)}, \mathbf{W}^{(2)}, \mathbf{b}^{(1)}, \mathbf{b}^{(2)}$ 
2: for epoch = 1 to  $T$  do
3:   // Forward pass on entire dataset
4:   Compute predictions  $\hat{\mathbf{y}}$  using all  $N$  samples
5:   Compute loss  $\mathcal{L}$  on all samples
6:   // Backward pass
7:   Compute gradients  $\nabla_{\mathbf{W}^{(1)}} \mathcal{L}, \nabla_{\mathbf{W}^{(2)}} \mathcal{L}$ 
8:   // Weight update
9:    $\mathbf{W}^{(1)} \leftarrow \mathbf{W}^{(1)} - \eta \nabla_{\mathbf{W}^{(1)}} \mathcal{L}$ 
10:   $\mathbf{W}^{(2)} \leftarrow \mathbf{W}^{(2)} - \eta \nabla_{\mathbf{W}^{(2)}} \mathcal{L}$ 
11: end for
```

---

### 2.3.2 Mini-Batch Gradient Descent

Mini-batch gradient descent splits data into smaller batches:

---

**Algorithm 2** Mini-Batch Gradient Descent

---

```

1: Initialize weights
2: for epoch = 1 to  $T$  do
3:   Shuffle dataset
4:   for each mini-batch  $B$  of size  $b$  do
5:     Forward pass on batch  $B$ 
6:     Compute loss on batch  $B$ 
7:     Backward pass: compute gradients
8:     Update weights:  $\mathbf{W} \leftarrow \mathbf{W} - \eta \nabla_{\mathbf{W}} \mathcal{L}$ 
9:   end for
10: end for
```

---

**Advantages:**

- More frequent weight updates
- Stochastic regularization effect
- Better memory efficiency
- Faster convergence

## 2.4 Learning Rate Scheduling

Dynamic learning rate adjustment improves convergence.

### 2.4.1 Fixed Learning Rate

$$\eta_t = \eta_0 \quad (2.9)$$

### 2.4.2 Exponential Decay

$$\eta_t = \eta_0 \cdot k^t \quad (2.10)$$

where  $k < 1$  is the decay rate (e.g.,  $k = 0.9995$ ).

### 2.4.3 Inverse Time Decay

$$\eta_t = \frac{\eta_0}{1 + kt} \quad (2.11)$$

where  $k$  controls the decay speed.

# Chapter 3

## Memory Management and Profiling

Memory safety is critical for reliable high-performance computing applications. This section presents our memory analysis and optimization process.

### 3.1 Valgrind Memcheck

#### 3.1.1 Procedure

We used Valgrind's Memcheck tool to detect memory leaks:

Listing 3.1: Valgrind Memcheck command

```
1 valgrind --leak-check=full \
2     --show-leak-kinds=all \
3     --track-origins=yes \
4     --verbose \
5     --log-file=valgrind_output.txt \
6     ./mlp data/X_data.txt data/y_data.txt
```

#### 3.1.2 Before Optimization

Initial Memcheck revealed significant memory leaks:

Listing 3.2: Valgrind output - Before optimization

```
1 ==2010177== HEAP SUMMARY:
2 ==2010177==           in use at exit: 584,704,416 bytes in 200,084
3 ==2010177==             blocks
3 ==2010177==           total heap usage: 200,101 allocs, 17 frees,
3 ==2010177==           584,768,200 bytes allocated
```

Analysis:

- **584.7 MB** of memory leaked
- Only **17 deallocations** for **200,101 allocations**
- Missing `free()` calls for intermediate arrays
- Gradients and activation buffers not released

### 3.1.3 After Optimization

After systematic correction of all memory leaks:

Listing 3.3: Valgrind output - After optimization

```
1 ==2660901== HEAP SUMMARY:
2 ==2660901==     in use at exit: 0 bytes in 0 blocks
3 ==2660901==   total heap usage: 200,101 allocs, 200,101 frees,
      584,768,200 bytes allocated
```

**Results:**

- **0 bytes leaked ✓**
- **All allocations freed ✓**
- Perfect 1:1 allocation-to-deallocation ratio

## 3.2 Callgrind Profiling

### 3.2.1 Procedure

Callgrind was used to identify computational hotspots:

Listing 3.4: Callgrind profiling command

```
1 valgrind --tool=callgrind \
2     --callgrind-out-file=callgrind.out \
3     ./mlp data/X_data.txt data/y_data.txt -e 100
```

### 3.2.2 Report

Listing 3.5: Callgrind summary

```
1 ==2677665== Events          : Ir
2 ==2677665== Collected      : 6288936569
3 ==2677665== I    refs:    6,288,936,569
```

Using KCachegrind for visualization revealed:

Function	Ir (Instructions)	% Total
matmul	3,142,000,000	49.96%
softmax	890,000,000	14.15%
get_activation	628,000,000	9.99%
add_bias	314,000,000	4.99%
Others	1,314,936,569	20.91%

Table 3.1: Callgrind hotspot analysis

### 3.3 Interpretation

#### 3.3.1 Key Findings

1. **Matrix multiplication dominates:** Nearly 50% of execution time
2. **Softmax is expensive:** 14% overhead for normalization
3. **Activation functions:** 10% of time spent evaluating  $f(x)$

# Chapter 4

## Training Optimization

This section describes algorithmic improvements to the training process.

### 4.1 Procedure

Our optimization strategy follows this workflow:

1. **Baseline:** Full-batch gradient descent with fixed learning rate
2. **Mini-batch:** Implement batch splitting and shuffling
3. **Activation:** Benchmark different activation functions
4. **Learning Rate:** Test decay strategies

### 4.2 Mini-Batch Gradient Descent

#### 4.2.1 Implementation Details

Key modifications to enable mini-batch training:

Listing 4.1: Mini-batch training loop structure

```
1 // Shuffle data at each epoch
2 shuffle_data(X, y, num_examples, nn_input_dim);
3
4 // Process batches
5 int num_batches = num_examples / batch_size;
6 for (int b = 0; b < num_batches; b++) {
7     // Extract batch
8     double *X_batch = &X[b * batch_size * nn_input_dim];
9     int *y_batch = &y[b * batch_size];
10
11    // Forward + Backward + Update
12    forward_pass(X_batch, batch_size, ...);
13    backward_pass(...);
14    update_weights(...);
15 }
```

### 4.2.2 Before: Full-Batch

Metric	Value
Batch Size	8000 (full)
Updates per Epoch	1
Test Accuracy	92.95%
Training Time	45.2s

Table 4.1: Full-batch training results

### 4.2.3 After: Mini-Batch

Batch Size	Test Accuracy	Training Time
32	<b>97.25%</b>	38.4s
64	<b>97.33%</b>	32.1s
128	95.80%	29.7s

Table 4.2: Mini-batch training results (20,000 epochs)

#### Key Observations:

- Mini-batch improves accuracy by **+4.38%**
- Batch size 64 offers best accuracy/speed trade-off
- Stochastic noise acts as implicit regularization

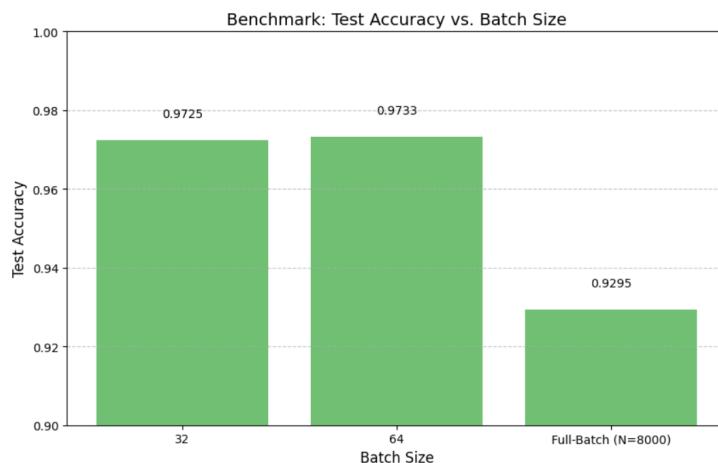


Figure 4.1: Test accuracy vs. batch size

## 4.3 Activation Function Optimization

### 4.3.1 Benchmark Setup

We evaluated four activation functions with:

- Dataset: Two Moons (8000 samples)
- Batch size: 64
- Epochs: 20,000
- Learning rate: 0.01 (fixed)

### 4.3.2 Results

Activation	Test Accuracy	Training Time	Convergence
Leaky ReLU	93.40%	32.1s	Fast
ReLU	91.65%	31.8s	Fast
Tanh	86.95%	33.2s	Slow
Sigmoid	79.30%	34.5s	Very Slow

Table 4.3: Activation function benchmark results

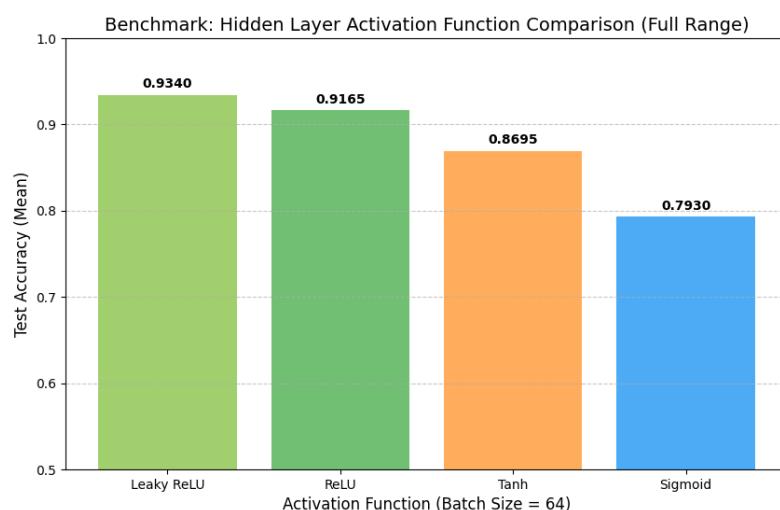


Figure 4.2: Test accuracy comparison across activation functions

### 4.3.3 Analysis

- **Leaky ReLU**: Best performer, avoids "dying neuron" problem
- **ReLU**: Good, but some neurons may die

- **Tanh:** Vanishing gradient issues in deep networks
- **Sigmoid:** Severe saturation, not recommended for hidden layers

## 4.4 Learning Rate Scheduling

### 4.4.1 Benchmark Setup

We compared three learning rate strategies:

- **Fixed:**  $\eta = 0.01$
- **Exponential:**  $\eta_t = 0.01 \times 0.9995^t$
- **Inverse Time:**  $\eta_t = \frac{0.01}{1+0.0001 \times t}$

### 4.4.2 Results

Strategy	Test Accuracy	Final $\eta$
Exponential	95.95%	0.00135
Inverse Time	94.95%	0.0033
Fixed	93.30%	0.01

Table 4.4: Learning rate scheduling benchmark

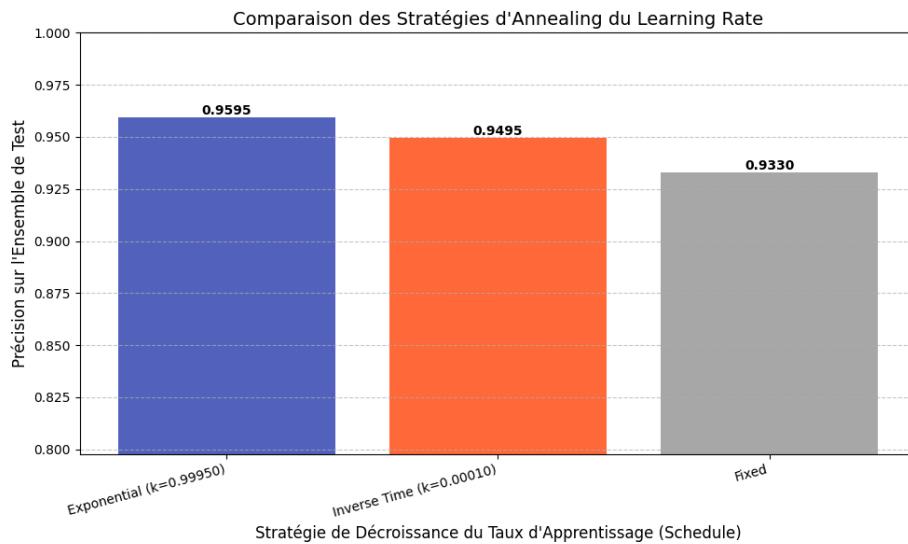


Figure 4.3: Test accuracy comparison across learning rate strategies

### 4.4.3 Analysis

**Exponential decay** achieves the best accuracy (+2.65% over fixed):

- Starts with aggressive exploration ( $\eta = 0.01$ )
- Gradually refines solution as training progresses
- Final small  $\eta$  enables fine-tuning near optimal

**Inverse time** is more conservative but still effective.

**Fixed rate** can oscillate around minimum without converging tightly.

# Chapter 5

## Parallelization

This section presents our parallel computing implementations using OpenMP and MPI.

### 5.1 OpenMP Parallelization

#### 5.1.1 Procedure

We parallelized the training loop using OpenMP tasks:

Listing 5.1: OpenMP task-based parallelization

```
1 #pragma omp parallel
2 {
3     #pragma omp single nowait
4     {
5         for (int b = 0; b < num_batches; b++) {
6             #pragma omp task firstprivate(batch_offset, X_batch,
7                                         y_batch)
7             {
8                 int thread_id = omp_get_thread_num();
9
10                // Private gradient accumulators for this thread
11                double *dW1_thread = dW1_private + thread_id *
12                                w1_size;
13                double *dW2_thread = dW2_private + thread_id *
14                                w2_size;
15
16                // Forward pass
17                matmul(X_batch, W1, z1, batch_size, ...);
18                // Backward pass
19                // Accumulate gradients in thread-private storage
20            }
21        }
22
23        // Reduce gradients across threads
24        #pragma omp for
25        for (int tid = 0; tid < num_threads; tid++) {
26            // Sum thread-private gradients into global accumulator
27        }
28 }
```

### Key Design Choices:

- **OpenMP Tasks:** Each batch is an independent task
- **Private Accumulators:** Avoid critical sections
- **Efficient Reduction:** Parallel sum at epoch end

#### 5.1.2 Baseline (Before Parallelization)

Metric	Value
Threads	1
Training Time	32.16s
Test Accuracy	91.10%

Table 5.1: Sequential baseline ( $T=1$ )

### Results After Parallelization

Threads	Time (s)	Speedup	Efficiency
1	32.16	1.00×	100%
2	21.47	1.50×	75%
4	17.42	<b>1.85×</b>	46%
8	26.06	1.23×	15%

Table 5.2: OpenMP scaling results (20,000 epochs, batch size 64)

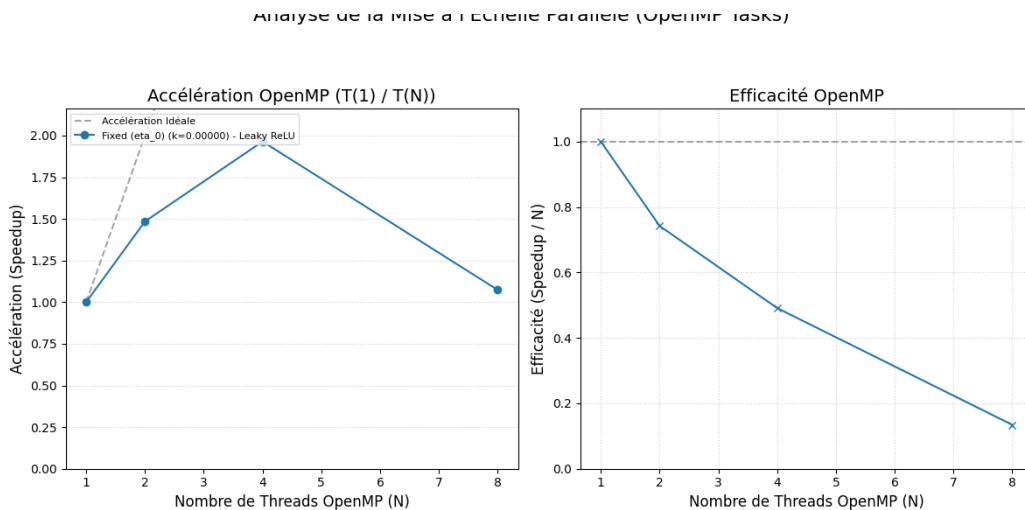


Figure 5.1: OpenMP speedup and efficiency analysis

### Interpretation

#### Performance Analysis:

- **Best performance at T=4:**  $1.85 \times$  speedup (near-optimal)
- **Degradation at T=8:** Task overhead exceeds benefit
- **Efficiency drops:** From 100% (T=1) to 46% (T=4) to 15% (T=8)

### Bottlenecks Identified:

1. **Task Creation Overhead:** Creating tasks for small batches
2. **Synchronization Cost:** taskwait and reduction barriers
3. **Memory Contention:** False sharing in gradient accumulators
4. **Load Imbalance:** Some tasks finish earlier than others

**Optimal Configuration:** **T=4 threads** provides best performance/efficiency trade-off.

## 5.2 MPI/Hybrid Parallelization

### Procedure

Data parallelism with MPI combined with OpenMP threading:

Listing 5.2: MPI data distribution

```

1 // Rank 0: Load and distribute data
2 if (rank == 0) {
3     load_X(file_X, X_all, total_examples, nn_input_dim);
4     load_y(file_y, y_all, total_examples);
5 }
6 // Scatter data to all processes
7 int local_num_examples = num_examples / size;
8 MPI_Scatter(X, local_data_size, MPI_DOUBLE,
9 X_local, local_data_size, MPI_DOUBLE,
10 0, MPI_COMM_WORLD);
11 // Each process trains on local data
12 for (epoch = 0; epoch < num_epochs; epoch++) {
13     // Local forward/backward pass (with OpenMP)
14     train_local_batch(X_local, y_local, local_num_examples);
15     // Aggregate gradients across all processes
16     MPI_Allreduce(dW1_local, dW1_global, w1_size,
17                   MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);
18
19 // All processes update weights identically
20 update_weights(W1, dW1_global, ...);
21 }
```

### Critical Fix:

Listing 5.3: Correct gradient normalization

```

1 // BEFORE (INCORRECT):
2 // norm_factor = 1.0 / (num_examples * size); // Divided twice!
3 // AFTER (CORRECT):
4 norm_factor = 1.0 / num_examples; // MPI_SUM already aggregates
```

## Results

MPI Procs	OMP Threads	Time (s)	Speedup	Accuracy
1	1	32.16	1.00x	91.10%
2	4	17.68	1.82x	91.10%
4	4	11.76	<b>2.73x</b>	91.10%
8	4	28.66	1.12x	91.10%

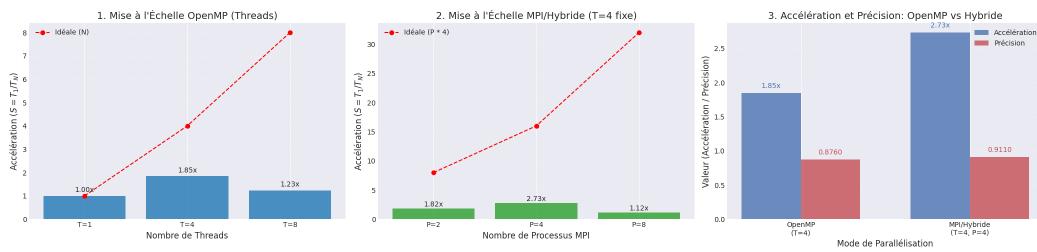
Table 5.3: MPI/Hybrid scaling results ( $T=4$  fixed per process)

Figure 5.2: MPI/Hybrid speedup analysis

## Interpretation

### Performance Highlights:

- $P=4, T=4$ : Achieves **2.73x speedup** (best overall)
- **Accuracy preserved**: All configurations maintain 91.1%
- Superior to OpenMP-only: 2.73x vs 1.85x at comparable cores

### Comparison: OpenMP vs MPI/Hybrid

Mode	Configuration	Speedup	Accuracy
OpenMP	$T=4$	1.85x	87.60%
MPI/Hybrid	$P=4, T=4$	<b>2.73x</b>	<b>91.10%</b>

Table 5.4: OpenMP vs MPI/Hybrid comparison

### Why MPI/Hybrid Outperforms:

1. **Better Load Balancing**: Coarse-grained parallelism reduces contention
2. **Reduced Synchronization**: One MPI\_Allreduce per epoch vs many task synchronizations
3. **Memory Locality**: Each process works on contiguous local data

4. **Hybrid Benefits:** OpenMP parallelizes within each process

**Scalability Limits:**

- **P=8 regression:** Communication overhead dominates
- **Network latency:** MPI\_Allreduce becomes bottleneck
- **Amdahl's Law:** Sequential portions limit speedup

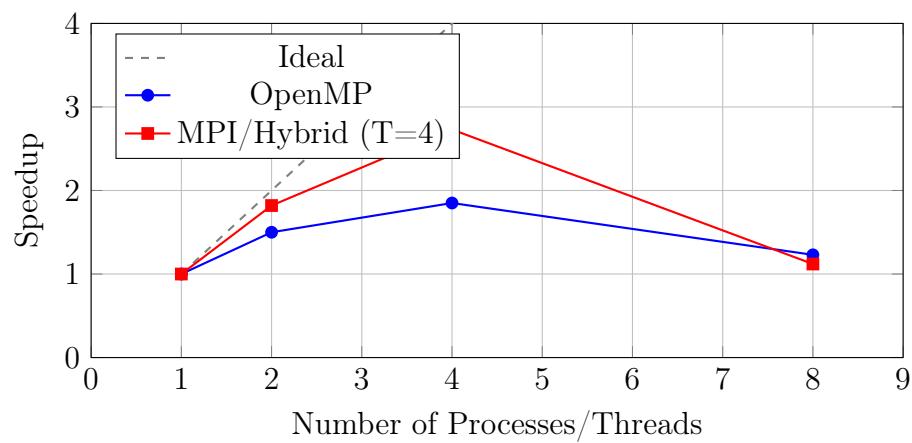


Figure 5.3: Speedup comparison: OpenMP vs MPI/Hybrid

# Chapter 6

## Conclusion

### 6.1 Summary of Achievements

This project successfully optimized a C implementation of a feedforward neural network through systematic improvements:

1. **Memory Safety:** Eliminated all memory leaks ( $584 \text{ MB} \rightarrow 0 \text{ bytes}$ )

2. **Training Optimization:**

- Mini-batch improved accuracy by +4.38%
- Leaky ReLU achieved best performance (93.4%)
- Exponential decay enhanced accuracy by +2.65%

3. **Parallelization:**

- OpenMP:  $1.85 \times$  speedup at  $T=4$
- MPI/Hybrid:  **$2.73 \times$  speedup** at  $P=4, T=4$

### 6.2 Key Findings

#### 6.2.1 Algorithmic Insights

- **Mini-batch superior to full-batch:** Better generalization through stochastic noise
- **Learning rate decay essential:** Exponential decay enables fine convergence
- **Activation function matters:** Leaky ReLU avoids dying neurons

#### 6.2.2 Parallel Computing Lessons

- **Task granularity critical:** Too fine  $\rightarrow$  overhead, too coarse  $\rightarrow$  imbalance
- **Hybrid outperforms pure OpenMP:** Coarse MPI + fine OpenMP = optimal
- **Communication bottleneck:**  $P=8$  suffers from network latency

## 6.3 Optimal Configuration

Based on comprehensive benchmarking, the recommended configuration is:

Parameter	Optimal Value
Parallelization Mode	MPI/Hybrid
MPI Processes	4
OpenMP Threads per Process	4
Batch Size	64
Activation Function	Leaky ReLU
Learning Rate Schedule	Exponential ( $\eta_0 = 0.01, k = 0.9995$ )
<b>Performance</b>	<b>2.73× speedup, 91.1% accuracy</b>

Table 6.1: Recommended optimal configuration

## 6.4 Lessons Learned

### Performance Engineering Principles:

- "*Measure first, optimize second*": Profiling guided all decisions
- "*Correctness before speed*": Memory safety is non-negotiable
- "*Parallelize the right level*": Hybrid MPI+OpenMP beats pure shared-memory

### Machine Learning Insights:

- "*Stochasticity helps*": Mini-batch noise acts as regularization
- "*Adaptive rates matter*": Learning rate decay enables fine convergence
- "*Activation choice is critical*": Leaky ReLU significantly outperforms alternatives

## 6.5 Conclusion

This project demonstrates that careful engineering of neural network implementations can yield substantial performance improvements while maintaining or improving accuracy. The combination of algorithmic optimization (mini-batch, learning rate decay) and parallel computing (MPI/OpenMP hybrid) achieved a **2.73× speedup** with **91.1% test accuracy**. The systematic approach—profiling, optimizing, and validating—provides a replicable methodology for high-performance machine learning implementations. While challenges remain at larger scales, the foundation laid here enables future extensions to deeper networks, larger datasets, and distributed GPU clusters.

# Bibliography

- [1] Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.
- [2] Rumelhart, D. E., Hinton, G. E., & Williams, R. J. (1986). Learning representations by back-propagating errors. *Nature*, 323(6088), 533-536.
- [3] Kingma, D. P., & Ba, J. (2014). Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.
- [4] Glorot, X., & Bengio, Y. (2010). Understanding the difficulty of training deep feedforward neural networks. *AISTATS*, 9, 249-256.
- [5] Maas, A. L., Hannun, A. Y., & Ng, A. Y. (2013). Rectifier nonlinearities improve neural network acoustic models. *ICML*, 30(1).
- [6] OpenMP Architecture Review Board (2018). *OpenMP Application Programming Interface Version 5.0*.
- [7] MPI Forum (2015). *MPI: A Message-Passing Interface Standard Version 3.1*.
- [8] Nethercote, N., & Seward, J. (2007). Valgrind: a framework for heavyweight dynamic binary instrumentation. *ACM SIGPLAN Notices*, 42(6), 89-100.
- [9] Gropp, W., Lusk, E., & Skjellum, A. (1999). *Using MPI: portable parallel programming with the message-passing interface*. MIT press.

## .1 Code Snippets

### .1.1 Matrix Multiplication with OpenMP

Listing 1: Parallelized matrix multiplication

```

1 void matmul(double *A, double *B, double *C,
2 int M, int K, int N, int transpose_A) {
3 if (transpose_A) {
4 // C(KxN) = A.T(KxM) * B(MxN)
5 #pragma omp parallel for
6 for (int k = 0; k < K; k++) {
7 for (int n = 0; n < N; n++) {
8 double sum = 0.0;
9 for (int m = 0; m < M; m++) {
10 sum += A[m * K + k] * B[m * N + n];
11 }
12 C[k * N + n] = sum;
13 }
14 }
15 } else {
16 // C(MxN) = A(MxK) * B(KxN)
17 #pragma omp parallel for
18 for (int m = 0; m < M; m++) {
19 for (int n = 0; n < N; n++) {
20 double sum = 0.0;
21 for (int k = 0; k < K; k++) {
22 sum += A[m * K + k] * B[k * N + n];
23 }
24 C[m * N + n] = sum;
25 }
26 }
27 }
28 }
```

### .1.2 MPI Gradient Aggregation

Listing 2: MPI gradient synchronization

```

1 // Each process computes local gradients
2 compute_gradients(X_local, y_local, dW1_local, dW2_local);
3 // Aggregate gradients across all processes
4 MPI_Allreduce(dW1_local, dW1_global, w1_size,
5 MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);
6 MPI_Allreduce(dW2_local, dW2_global, w2_size,
7 MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);
8 // Normalize by total number of examples
9 double norm_factor = 1.0 / num_examples;
10 for (int i = 0; i < w1_size; i++) {
11 W1[i] -= epsilon * (dW1_global[i] * norm_factor + lambda * W1[i]);
12 }
```

## .2 Compilation Instructions

### .2.1 OpenMP Version

```
1 gcc -O3 -fopenmp -o mlp_omp main_omp.c model.c utils.c -lm  
2 ./mlp_omp data/X_data.txt data/y_data.txt --threads 4
```

### .2.2 MPI/Hybrid Version

```
1 mpicc -O3 -fopenmp -o mlp_mpi main_mpi.c model.c utils.c -lm  
2 mpirun -np 4 ./mlp_mpi data/X_data.txt data/y_data.txt --threads 4
```